A thorough evaluation of the Language Environment Analysis (LENATM) system

many[1]

1

Author Note

Correspondence concerning this article should be addressed to many, . E-mail:

Abstract

In the previous decade, dozens of studies involving thousands of children across several research disciplines have made use of a combined daylong audio-recorder and automated algorithmic analysis called the LENA system, which aims to assess children's language environment. However, validation of the system lags behind its growing importance in the research domain. Here we assess LENA accuracy across its key outcome measures: speaker classification, adult word counts, child vocalization counts, and conversational turn counts. Our assessment is based on manual (LENA-output naive) random and period sampling, in (a) populations similar to LENAs original training data from North American English-learning children, (b) another dialect of English (UK), and (c) in a different language and socio-cultural setting (rural Bolivia). We find reasonably high accuracy across some measures, with more problematic levels of performance across others. We find little difference in accuracy as a function of dialect, language, or socio-cultural setting. Whether LENA results are "good enough" for a given research, educational, or clinical application depends largely on the specifics at hand; we conclude with a set of recommendations to help researchers make this determination for their goals.

A thorough evaluation of the Language Environment Analysis (LENATM) system

While nearly all humans eventually become competent users of their language(s), documenting the experiential context of early acquisition is crucial for both theoretical and applied reasons. Regarding theory, there are many open questions about what kinds of experiences and interactions are necessary, sufficient, or optimal for supporting language development. Moreover, the ability to accurately and quickly assess an infant's state of development at a given point in time is of central importance for clinical purposes, both for children with known risks of language delays and disorders, and those who might not be identified based on risk factors. Reliable assessments are also crucial for measuring intervention efficacy.

One approach that has been making its way into the mainstream literature across applied and basic research on language and cognition relies on day-long recordings gathered with LENA(R) products (e.g. Greenwood et al, 2011; Oller et al 20 2010; Gilkerson et al 2017, inter alia), analyzed by LENA's automated, closed-source algorithms. As we summarize below, this approach has many advantages, which may explain its expanding popularity. While dozens of papers over the past decade have used LENA's automated output, only a handful include validity estimates [e.g. Weisleder & Fenald 2013; Zimmerman et al 2009, d'Apice et al 2019], even fewer where validity estimation was the primary focus of the paper (e.g. Ganek & Eriks-Brophy, Bulgarelli & Bergelson, in press @ BRM, Lehet Dilley Houston; for a fuller report, see Cristia et al, under review). As a result, few studies report sufficient details about validation accuracy for one or more metrics, limiting systematic review or meta-analytic assessment (Cristia et al., under review).

The work undertaken thus far also has some limitations, which are described further in the "Previous Validation" section below, and mentioned briefly next. First, many validations or evaluations of LENA omit analysis of the less-directly "relevant" categories of input like noise, silence, or overlap. Second, many LENA evaluations rely on LENA's own output as a starting point to either select portions of the file for manual annotation, or use its

segmentation into talker-turns as the unit of analysis (without conducting independent segmentation). Both decisions can lead to inflated accuracy estimates. Here we endeavor to conduct an evaluation that is fully independent of the LENA algorithms automated output assessment, permitting a systematic, extensive, and independent evaluation of the key automated metrics provided by this system. Summarily, we report the LENA(R) algorithms' performance based on random or periodic sampling from a set of recordings, which (as detailed and argued below) is preferable to other types of sampling. We conduct our analysis on (a) a sample of children similar to the LENA(R) training set (i.e. infants and toddlers, growing up in North American English-speaking homes, [LTR-06-2]) (b) check for generalization to a different dialect (UK English); and (c) extend our analysis to a different language and setting (Tsimane'-speaking homes in rural Bolivia).

Brief introduction to LENA(R) products. The LENA(R) system consists of a hardware component (a lightweight, sturdy, and easy-to-use recording device worn by a child in specialized clothing) and a suite of proprietary computer programs designed to provide automated quantitative analyses of the auditory environment and the child's own vocalizations.

The LENA software was developed over an extensive corpus of full day audio recordings from a child's perspective using their patented recording hardware [Xu et al, 2009]. The original dataset included over 65,000 hours of recording across 329 American English-speaking families chosen for diversity in child age (1-42 months) and socio-economic status [LTR-6-2]. From these recordings, half-hour selections from 309 recordings were transcribed and annotated for the purpose of developing the algorithm, and an additional 60 minutes from each of 70 recordings were transcribed and annotated for testing the result [ibid].

The resulting LENA software takes each audio recording and processes it incrementally in short windows, extracting a variety of acoustic features which are used to classify the audio stream into segments of at least 600 ms in length (or longer for some of the categories)

using a Minimum Duration Gaussian Mixture Model (MDGMM, [Xu et al, 2008]). Silence may be included to "pad" segments to this minimum duration. The segments are classified according to a set of broad speaker classifications - Male Adult, Female Adult, "Key" Child (i.e. the one wearing the recorder) and Other Child - and non-speaker classifications - Noise, Television (including any electronics), Overlap (speech overlapped with other speech or nonspeech sounds), and Silence (SIL). With the exception of Silence, these classifications are then passed through a further likelihood test and labeled as "Near" (high probability) or "Far" (low probability), yielding the following classes: MAN/MAF, FAN/FAF, CHN/CHF, CXN/CXF, NON/NOF, TVN/TVF, OLN/OLF. Given the large number of acronyms and labels of various kinds, we provide a listing of relevant LENA abbreviations on Table xx; in what follows, we highlight the specific labels that we evaluate in the present work in bold.

After this broad speaker classification step, Female or Male Adult "Near" segments (FAN and MAN) are further processed using an adaptation of the Sphinx Phone Decoder [Lamere et al, 2003] in order to form an automated estimate of the number of words in each segment (adult word count, or **AWC**). Key Child (CHN) segments are further processed to sub-classify into vegetative noises (VEG), crying (CRY), and speech-like vocalization (VOC). LENA provides the count (child vocalization count, or CVC) and duration of these VOC sub-segments. A further metric, Conversational turn counts (CTC), reflects the number of alternations between an adult and the key child (or vice versa), bounded by 5 s of non-speech.

Table xx. A partial listing of common LENA abbreviations and their meanings.

Previous validation work.   A recent systematic review (Cristia, Bulgarelli, & Bergelson, submitted) found 23 papers, containing 28 studies, reporting on the accuracy of LENA's labels and derived metrics (AWC, CVC, CTC). They conclude there are "reasonably good results [overall]: over 61% for recall and precision based on 11-12 non-independent studies; correlations for AWC mean r=.79, on n=11, with a mean RER[Relative Error Rate]=10% on n=11; CVC mean r=.76, n=5, with a mean RER=1% on n=5. The exception

to this general trend towards good performance was CTC, with a mean r=.31, n=5, RER=-64% on n=2."

While this systematic review sought to take stock of the current literature using LENA reportings, there were several limitations that make it likely that these measures do not provide a full or fair picture of LENA output accuracy. First, for the majority of included studies, the validity report was not fully evaluated by peer review. Even if the study may have appeared in a peer-reviewed report, the LENA validation was often a secondary goal to support a different research objective, and therefore often lacked methodological details or even full results. For instance, Seidl et al. report on validation of LENA labels among children at familial risk for autism in a short appendix to the journal; in this Appendix, only confusions between female adult and child are mentioned – suggesting that confusions between key child and any other category (other child, male adult, silence, etc.) were ignored rather than counted as an error. While this approach may be reasonable for a given study's goals, it has the undesirable side effect of reducing the granularity of our understanding. Second, previous studies typically did not take silence, noise, or overlap into account in the reported confusion matrices or other accuracy measures, particularly within segments. That is, if a LENA segment labeled "key child" contained one second of silence and two seconds of speech, the full three second clip was often tagged as "correct" though it was only 67% correct, leading to an overestimation of the accuracy of the "key child" label.

Third, a majority of previous validation studies used the LENA output itself to select the sections that would be annotated for validation (in Cristia et al., subm, this held for 14/25 studies that specified the method of selection). For instance, clips may have been selected for manual annotation on the basis of high AWC and/or CTC according to the algorithm. This unfortunately leads to biased sampling: Since LENA only counts words within FAN and MAN segments and conversational turns involving FAN/MAN alternations with CHN in close proximity, high AWC or CTC can only occur in sections of the recording that are "clean" enough for the algorithm to parse; otherwise, most of the section would have

been classified as overlap (OLN), which does not count towards AWC or CTC. This would tend to bias these reports toward a higher level of accuracy than would be obtained overall across the full recording.

Forth, previous validation work has typically focused on a single corpus, participant population, age range, and language. As a result, it is difficult to assess whether a numerical difference in accuracy found across two papers is statistically significant. Even if this difference is large, it is challenging to decide whether this is due to a difference in the way the corpus was constituted and annotated, on how LENA fares with that population, age range, and language.

The present work. We sought to assess the validity of LENA's metrics through an approach that complements the preceding literature (including our own work). Specifically, we report an evaluation of all speech labels and some non-speech labels of LENA (silence (SIL), overlap near (OLN), overlap far (OLF)), as well as LENA's key derived metrics: adult word count (AWC), conversational turn count (CTC), and child vocalization count (CVC). We aim to address several of the potential limitations highlighted above.

First, to maximally avoid potential bias in our annotations, we used random or periodic sampling (detailed below) to choose which sections of daylong recordings to annotate, and did not give annotaters access to the LENA output. Second, to allow assessment of the accuracy of the segmentation itself as well as categorical labeling, evaluation is done at the level of 10 ms frames. This will allow us to capture a much finer-grained representation of the auditory environment (i.e. if LENA classified a 2 s audio segment as FAN, but .8 s of this was actually non-speech noise, LENA would be credited only for the proportion that was correct).

Third, to gain traction on generalizability rather than focusing on a single sample that either mirrors or diverges from LENAs original population, we include five corpora. Three corpora based on the same population, language, and dialect LENA was established with; a fourth corpus that allows an extension to a different dialect of English; and a fifth corpus

that is an extension to a totally different recording condition (a rural setting, with large families and many children present, in a different language). The age range also varies a great deal, particularly in this last corpus.

Finally, the present study relies on a collaborative annotation effort across several labs. Critically, this means that we can "fairly" compare annotations across our corpora: annotation decisions were either identical or comparable, and all analyses were identical. This allows us to more readily answer questions regarding differences in reliability as a function of e.g. child age and language.

Methods

Corpora.   The data for the evaluation comes from two different projects. The largest one is the ACLEW project (Bergelson et al., 2017b Databrary volume; Soderstrom et al AMPPS under review); in this paper we focus on four different corpora of child daylong recordings that have been pooled together, sampled, and annotated in a coordinated manner. These four corpora are: the Bergelson corpus ("BER") from US English families from the upstate New York area (Bergelson, 2016), the LuCiD Language 0–5 corpus ("L05") consisting of English-speaking families from Northwest England (Rowland et al., 2018), the McDivitt and Winnipeg corpora ("MCD") of Canadian English families (McDivitt & Soderstrom, 2016), and the Warlaumont corpus ("WAR") of US English from Merced, California (Warlaumont et al., 2016). Some recordings in BER, and all recordings in MCD and WAR are available from HomeBank repository (VanDam et al., 2016). The second project contains a single corpus collected from Tsimane' speaking families in Bolivia (Scaff et al., in prep.; here "TSI").

Key properties of the five corpora are summarized in Table 2. Each corpus consists of daylong (4–16 hour) at-home recordings; each corpus samples from a unique community. Corpora span languages and dialects; socioeconomic environment varying both within and across corpora. In each recording, the key child ("participant") wears a LENA recorder in a

special vest throughout a normal day.

For the four ACLEW corpora, out of the 106 recorded participants, daylong recordings from 10 infants from each corpus were chosen for manual annotation, selected to represent a diversity of ages (0–36 months) and socio-economic contexts. In the SOD corpus, sensitive information was found in one of the files, and thus one child needed to be excluded. The tenth day for this corpus was a second day by one of the 9 included children. From each daylong file, fifteen 2-minute non-overlapping audio (with a 5-minute context window) were randomly sampled from the entire daylong timeline for manual annotation. This 30 minute sample corresponded to approximately 1 minute of annotated speech per key child (collapsing across all speaker categories).

The TSI corpus consisted of 1-2 daylong recordings from 12 infants, out of the 25 recorded from field work that year; the other 13 had been recorded using other devices (not the LENA DLP). From these daylong files, 1-minute segments were sampled in a period fashion. That is, for each day, we skipped the first 33 minutes to allow the family to acclimate to the recorder, and then extracted 1 minute (with a 5-minute context window) every 60 minutes, until the end of the recording was reached. This resulted in 4 to 16 minutes of manually annotated audio per child per recording (mean = 12.64 minutes), and an average of 3 minutes of speech per key child (collapsing across all speaker categories).

We chose to sample 1 or 2 minutes at a time (Tsimane, and ACLEW corpora, respectively) because conversations are likely to be bursty (Goh & Barbasi; cf. Slone et al Fausey et al ICIS18 session). That is, it is likely the case that speech is not produced at a periodic rate (e.g., one phrase every 20 seconds), but rather it occurs in bursts (a conversation is followed by a long period of silence between the conversational partners, followed by another bout of conversation, perhaps with different interlocutors, followed by silence, and so on). In this context, imagine that you sample a 5-second stretch. If you find speech in that stretch, then it is likely you have by chance fallen on a conversation bout; if you do not find speech, then you have likely found a silence bout. If you were to extend that

selection out to several minutes, then it is likely that you will simply add more material from the same type (i.e. conversation bout or silence bout). As a result, any sampling method that favors medium-sized stretches (5-15 minutes) will tend to end up with samples that are internally homogeneous (throughout the 5 minutes, there is a conversation), but highly heterogeneous as a collection if sampling is random throughout the day. This in turn would lead to artificially high correlations between LENA and human metrics in all dependent variables (e.g., since probably little speech, fewer turns, and fewer child vocalizations will be found in silent bouts and more in conversational ones). Thus, our strategy of sampling randomly/periodically and in short stretches is more likely to represent both conversational and short bouts, and to capture finer-grained variation in speech quantity.

In the 5 corpora, the 1- or 2-min samples were annotated for all hearable utterance boundaries and talker ID. In ACLEW corpora [1], talker IDs reflected unique individual talkers, but were coded in such a way to readily allow mapping onto LENAs talker categories (e.g. key child, other children, female adult, male adult). In the TSI corpus, only the key child and one female adult whose voice recurred throughout the day were individually identified, with all other talkers being classified on the basis of broad age and sex into male adult, female adult, and other children. The ACLEW datasets had other coding levels which will not be discussed here.

Processing.   Several different time units are needed to clarify how different metrics are calculated (see Figure XX). Clips refer to the 1- or 2-minute samples extracted from recordings. This is the basic unit at which child vocalization counts and conversational turn counts can be established. In addition, since most previous work evaluating adult word counts did so at the clip level, we do so here as well.

The other metrics require a more detailed explanation, conveyed graphically in Figure XX. The stretch of time that has been assigned to a speech or non-speech category by LENA

---

[1]see Casillas et al., 2017a, 2017b for the general annotation protocol, and Soderstrom et al., submitted to AMPPS, for an introduction to the databases

Figure 1. Levels at which performance is evaluated.

is a segment. In one clip, there may be just one long segment (e.g., the whole clip has been assigned to Silence by LENA); or there may be more (e.g., the first 5 seconds are attributed to the key child, then there is a 50-second Silence segment, and the final 5 seconds are attributed to a Female Adult). In LENA's automated analysis, only one of these categories may be active at a given point in time. In contrast, we typically speak of utterances or vocalizations to refer to stretches of speech detected by humans and assigned to different talkers. Again, clips may have zero or more utterances. Unlike LENA, however, a given point in time may be associated with multiple speakers. Given that there need not be a one-to-one correspondence between LENA segments and human utterances, we need to define smaller time units that can be used to check for agreement. In this paper, we use 10 ms frames. This is the basic time unit used for all classification accuracy estimations, which are introduced in more detail in the next subsection.

LENA classification accuracy.   Our first goal was to establish LENA talker tag accuracy, particularly for the four broad LENA talker categories (key child, other child, female adult, male adult; or CHN, CXN, FAN, MAN), but taking into account other

categories (with some limitation on their interpretation). We calculated accuracy in two complementary ways. First, we used three frame-based, standard metrics of speech and talker segmentation to allow direct comparison with other systems in the speech technology literature (False Alarms, Misses, Confusion Rate). We also use Diarization Error Rate, which is derived from the first three metrics; together these provide a stringent and standard test of accuracy. Second, we used frame-based precision and recall of each category to provide an intuitive representation of the error patterns shown by this system. In both cases, electronic voices were rare in the ACLEW annotations and not coded at all (by design) in the Tsimane data, they were mapped onto silence.

Speech and talker segmentation metrics. The original coding was converted using custom-written python scripts into rttm format (REF), a text-based format indicating, for each vocalization, its start time, duration, and speaker. This representation was used in pyannote.metrics (REF) to compute four standard diarization metrics: rate of false alarm for speech, rate of misses for speech, rate of confusion between talkers, and the derived diarization error rate (DER). These are calculated with the following formulas at the level of each clip, where FA (false alarm) is the number of frames during which there is no talk according to the human annotator but during which LENA found some talk; M (miss) is the number of frames during which there is talk according to the human annotator but during which LENA found no talk; C (confusion) is the number of frames correctly classified as LENA as containing talk, but whose voice type has not been correctly identified (when the LENA model recognizes female adult speech where there is male adult speech for instance), and T is the total number of frames that contain talk according to the human annotation:

- false alarm rate = FA/T,
- miss rate = M/T,
- confusion rate = C/T,
- DER = (FA+M+C)/T,

In the human annotation, there is no class representing overlapping speech as such. For

the sake of completeness and full comparison with the LENA model, time where two or more different speech sources were active at the same time, according to the human annotators, have been mapped to the class "overlap" post hoc. This allows us to compare this Overlap class to LENA's OLN (and, for the precision/recall analysis introduced next, OLF) by the LENA model. Therefore, the confusion rate is computed based on the matches in Table XX.

Table XX: Correspondances between LENA and our human annotation. *Electronic voices were only annotated in the ACLEW dataset. Although some Tsimane' families listen to the radio, radio speech was not annotated in the TSI corpus.

Please remember that the overlap category is not defined the same way as the LENA overlap category. For LENA, overlap between any two categories falls within overlap – i.e., CHN+TV would be counted towards overlap as would FAN+FAN; whereas for us, only overlap between two talker categories (e.g., key child and female adult) counts as overlap. Similarly, the TVN is not equivalent to the electronic speech tag in the ACLEW coding.

Precision and recall.   This evaluation looks in more detail at the pattern of errors, by assessing how LENA and human annotators agreed and disagreed. In both precision and recall, the numerator is the intersection between a LENA tag and a human tag (e.g., the number of frames that LENA classified as CHN and the annotator classified as Key child; notice, there is no constraint that these two categories by conceptually the same). The denominator differs: To calculate precision, we divide that number by the total number of frames attributed to a category by LENA, whereas for recall, we divide by the total number of frames attributed to a category by the human annotator.

CVC and CTC evaluation.   From the human annotation, each vocalization by the key child counted towards the total Child Vocalization Count (CVC) for a given clip. For the Conversational Turn Count (CTC), a sequence of key child and any adult (or vice versa) within 5 seconds counted towards the clip total CTC.

AWC evaluation.   For the AWC portion of this evaluation, we could only use transcriptions from the four ACLEW corpora, since the TSI corpus has not been transcribed

(and thus lacks word counts). In contrast, annotators for the four ACLEW corpora were proficient in the language spoken in the daylong recording, and transcribed all adult speech based using canonical lexical forms (e.g. "wanna", not "want to") in keeping with minCHAT format (MacWhinney).

Reference adult word counts were determined by counting all unambiguously transcribed words spoken by adult talkers. This was achieved by first discarding all non-lexical transcript entries such as non-linguistic communicative sounds, paralinguistic markers, and markers indicating incomprehensible speech. In addition, all utterances from the key child and other children were omitted from the Adult Word Count. The remaining orthographic entries separated by whitespaces were then counted as gold standard target words for LENA to detect.

As for LENA word counts, the regions sampled for manual annotation were not guaranteed to perfectly align with LENA utterances. Of all LENA utterances overlapping with the annotated speech, 14% had only partial overlap with the annotations. To match LENA AWCs with the annotated word counts, words from the partially overlapping LENA utterances were included in proportion to the amount of overlap between the LENA turn and the reference segment in question (e.g., if 10% of a LENA-estimated utterance was overlapping with a manually-annotated utterance, 10% of the total LENA AWC estimate for the given turn was included in the LENA word count estimate for that reference segment).

Results

Before starting, we provide some general observations based on the human annotation. Silence is extremely common, constituting 79% of the frames. In fact, 45% of clips contained no speech by any of the human speaker types (according to the human annotators). As for speakers, female adults make up 11% of the frames, the child contributes to 4% of the frames, whereas male adult voices, other child voices, and electronic voices are found in only 1% of the frames each. Overlap makes up the remaining 3% of the frames. The following

consequences ensue: if frame-based accuracy is sought, a system that classifies every frame as silence would be 79% correct. This is of course not what we want, but it indicates that systems adapted to this kind of speech should tend to have low "false alarm" rates, i.e. a preference for being very conservative as to when there is speech. If the system does say there is speech, then it had better say that this speech comes from female adults, who provide a great majority of the speech. In second place, it should be key child. Given that male adults and other children are rare, a system that makes a lot of mistakes in these categories may still have a good global performance, because these categories are extremely rare.

LENA classification accuracy: False alarms, misses, confusion. Our first analysis is based on standard speech technology metrics, which put errors in the perspective of how much speech there is. That is, if 10 frames are wrong in a file where there are 100 frames with speech, this is a much smaller problem than if 10 frames are wrong in a file where there is 1 frame with speech. In other words, these metrics should be considered relative error metrics. One problem, however, emerges when there is no speech whatsoever in a given file. In the speech technology literature, this is never discussed, because most researchers working on this are basing their analyses on files that have been selected to contain speech (e.g., recorded in a meeting, or during a phone conversation). We still wanted to take into account clips with no speech inside because it is key for our research goals: We need systems that can deal well with long stretches of silence, because we want to measure how much speech children hear. Indeed, as mentioned above, 45% of our clips had no speech whatsoever. In these cases, the false alarm, miss, and confusion rates are all undefined, because the denominator is zero. In all likelihood, this leads to an overestimation of LENA's performance, because potential false alarms in these files are not counted against the system. It also occurred that there was just a little speech; in this case, the denominator is very small, and therefore the ratio for these two metrics ended up being a very large number. To avoid such outliers having an undue impact on our report, we present medians (rather than means).

There were, a priori, several ways of analyzing the data:

- collapsing near and far together (i.e., CHN and CHF were mapped onto a single CH category)

- treating the near and far categories separately (i.e., CHN and CHF are both treated as "speakers", but not the same one)

- not considering TV as a speaker category, since it is conceptually not identical to the electronic voices detected by ACLEW human annotators; in this case, the gold annotations should also map electronic voices to non-speech or silence

- not considering OLN as a speaker category, since it is not conceptually identical to the overlap derived from humans' annotating different speaker categories.

We thought the most informative decision would be to report on several of these settings, albeit briefly. We start with the situation that yields the best LENA performance: Electronic voices in the gold annotation are mapped onto silence, so that the categories found in the human annotation are FEM, MAL, CHI, OCH, and overlap; in the LENA annotation, only CHN, FAN, MAN, and CXN are considered speakers (with all far categories, TVN, and OLN all mapped onto silence). In this setting, LENA's false alarm (i.e., saying that someone was speaking when they were not) had a median of 12%, whereas the miss rate had a median of 49%. The confusion rate, as mentioned above, is only calculated for the correctly detected speech (i.e., not the speech that was missed, which counts towards the miss rate, nor the speech that was falsely identified, which is considered in the false alarm). The confusion rate was very low, with a median of 10%. These three metrics can be added together into a single "diarization error rate". The median diarization error rate over the clips that had some speech was 79%.

If electronic voices in the gold annotation are still mapped onto silence and in the LENA annotation, CHN, FAN, MAN, CXN as well as OLN are considered speakers (with all far categories as well as TVN mapped onto silence), so that the human categories considered were CHI, FEM, MAL, OCH, and overlap; and the LENA categories considered were CHN,

FAN, MAN, CXN, and OLN. In this setting, LENA's false alarm, missed, and confusion rate medians were 33%, 21%, and 28% respectively, for a total median diarization error of 82%. Performance likely degrades because OLN is not picking up the same regions as the overlapping speech found in the human annotations.

Next, we allowed the electronic voices segmented by humans, and TVN among the LENA speaker categories, to be considered during the evaluation (rather than mapping them all to non-speech or silence), so that the human categories considered were CHI, FEM, MAL, OCH, overlap, and electronic; and the LENA categories considered were CHN, FAN, MAN, CXN, OLN, and TVN. LENA's false alarm, missed, and confusion rate medians were 40%, 18%, and 30% respectively, for a total median diarization error of 88%. Performance likely degrades because TVN is not picking up the electronic speech segmented by ACLEW annotators.

Finally, we declared the maximum possible number of categories: The human categories considered were still CHI, FEM, MAL, OCH, overlap, and electronic; but the LENA categories considered were CHN, FAN, MAN, CXN, OLN, TVN, CHF, FAF, MAF, CXF, OLF, TVF. LENA's false alarm, missed, and confusion rate medians were 74%, 7%, and 41% respectively, for a total median diarization error of 122%. Performance degrades because everything is treated as speech, leading to huge apparent false alarm rates.

LENA classification accuracy: Precision and recall. By now, we have established that the best performance (when "far" labels such as CHF and OLF are mapped onto silence, as are TVN and OLN), the overall relative diarization error rate is about 79%, due mainly to missing speech (49%), with false alarms (12%) and confusion between talker categories (10%) constituting a relatively small proportion of errors. However, this metric may not capture what our readers are interested in, for two reasons. First, this metric gives more importance to correctly classifying segments as speech versus non-speech (False alarms + misses) than confusing talkers (confusion). Second, many LENA adopters use the system not to make decisions on the sections labeled as non-speech, but rather on sections labeled as

speech, and particularly those labeled adults and key child. The metrics above do not give more importance to these two categories, and do not give us insight on the patterns of error made by the system. Looking at precision of speech categories is crucial for users who interpret LENA's estimated quantity of adult speech or key child speech, as low precision means that some of what LENA called e.g. key child was not in fact the key child, and thus it is providing overestimates. Looking at recall may be most interesting for adopters who intend to employ LENA as a first-pass annotation: the lower the recall, the more is missed by the system and thus cannot be retrieved (because the system labeled it as something else, which will not be inspected given the original filter). Recall also impacts quantity estimates, since it indicates how much was missed of that category.

Therefore, this subsection shows confusion matrices, containing information on precision and recall, for each key category. For this analysis, we collapsed over all human annotations that contained overlap between two speakers into a category called "overlap". Please remember that this category is not defined the same way as the LENA overlap category. For LENA, overlap between any two categories falls within overlap – i.e., CHN+TV would be counted towards overlap; whereas for us, only overlap between two talker categories (e.g., key child and female adult) counts as overlap.

We start by explaining how to interpret one cell in Figure (precision): Focus on the crossing of the human category FEM and the LENA category FAN; when LENA tags a given frame as FAN, this corresponds to a frame tagged as being a female adult by the human 52% of the time. This category, as mentioned above, is the most common speaker category in the audio, so that over 57k frames (representing 52% of the frames tagged as FAN by LENA) were tagged as being female adult by both the human and LENA. The remaining 2, 0, 1, 8, 0, and 37% of frames that LENA tagged as FAN were actually other categories according to our human coders: 37% were silence, 8% were in regions of overlap between speakers or between a speaker and an electronic voice, and 3% were due to confusions with other speaker tags. Inspection of the rest of the confusion matrix shows that,

Figure 2

other than silence, this is the most precise LENA tag.

Precision for CHN comes in secondplace, at 39%; thus, fewer than half of the frames labeled as being the key child are, in fact, the key child. The majority of the frames, LENA incorrectly tagged as being the key child are actually silence (or rather, lack of speech) according to the human annotator (44%), with the remaining errors being due to confusion with other categories: About 8% of them are actually a female adult; 2% are another child; and 7% are regions of overlap across speakers, according to our human coders.

MAN and CXN score similarly, 8 and 6% respectively, meaning that less than a tenth of the areas LENA tagged as being these speakers actually correspond to them. As with the key child, most errors are due to LENA tagging silent frames as these categories. However, in this case confusion with other speaker tags is far from negligible. In fact, the most common speaker tag in the human annotation among the regions that LENA tagged as

being MAN were actually female adult speech (27%); and, for CXN, it was not uncommon to find a CXN tag for a frame human listeners identified as a female adult (17%) or the key child (6%). In a nutshell, this suggests extreme caution before undertaking any analyses that rely on the precision of MAN and CXN, since most of what is being tagged as such is silence or other speakers.

Another observation is that the "far" tags of the speaker categories do tend to more frequently correspond to what humans tagged as silence (74%) than the "near" tags (54%), and thus it is reasonable to exclude them from consideration. The relatively high proportion of near LENA tags that correspond to regions that humans labeled as silence could be partially due to the fact that the LENA system, in order to process a daylong recording quickly, does not make judgments on small frames independently, but rather imposes a minimum duration for all speaker categories, padding with silence in order to achieve it. Thus, any key child utterance that is shorter than .6 secs will contain as much silence as needed to achieve this minimum (and more for the other talker categories). Our system of annotation, whereby human annotators had no access whatsoever to the LENA tags, puts us in an ideal situation to assess the impact of this design decision, because any annotation that starts from the LENA segmentation should bias the human annotator to ignore such short interstitial silences to a greater extent than if they have no access to their tags whatsoever.

These analyses shed light on the extent to which we can trust the LENA tags to contain what the name indicates. We now move on to recall, which indicates a complementary perspective: how much of the original annotations were captured by LENA.

Again, we start with an example to facilitate the interpretation of this figure: The best performance for a talker category this time is CHN: Nearly half of the original frames humans tagged as being uttered by the key child were captured by the LENA under the CHN tag. Among the remaining regions that humans labeled as being the key child, 22% was captured by LENA's CXN category and 20% by its OLN tag, with the remainder spread out across several categories. This result can be taken to suggest that an analysis pipeline that uses the

Figure 3

LENA system to capture the key child's vocalizations by extracting only CHN regions will get nearly half of the key child's speech. Where additional human vetting is occuring in the pipeline, such researchers may consider additionally pulling out segments labeled as CXN, since this category actually contains a further 22% of the key child's speech. Moreover, as we saw above, over a third of these LENA tags corresponds to the key child, which means that human coders who are re-coding these regions could filter out the two thirds that do not.

Many colleagues also use the LENA as a first pass to capture female adult speech via their FAN label. Only 30 of the female adult speech can be captured this way. Unlike the case of the key child, missed female speech is classified into many of the other categories, and thus there may not exist an easy solution (i.e., one would have to pull out all examples of many other categories to get at least half of the original female adult). However, if the hope is to capture as much of the female speech as possible, perhaps a solution may be to also pull

out OLN regions, since these capture a further 25% of the original female adult speech and, out of the OLN tags, 16% are indeed female adults (meaning that human annotators re-coding these regions need to filter out 4 out of 5 clips, on average).

For the remaining two near speaker labels (MAN, CXN), recall averaged 15%, meaning that less than a quarter of male adult and other child speech is being captured by LENA. In fact, most of these speakers' contributions are being tagged by the LENA as OLN (mean across MAN and CXN 34%) or silence (mean across MAN and CXN 7%), although the remaining sizable proportion of misses is actually distributed across many categories.

Finally, as with precision, the "far" categories show worse performance than the "near" ones. It is always the case that a higher percentage of frames is "captured" by the near rather than the far labels. For instance, out of all frames attributed to the key child by the human annotator, 42% were picked up by the LENA CHN label versus 0% by the LENA CHF label. This result can be used to argue why, when sampling LENA daylong files using the LENA software, users need not take into account the "F" categories.

Child Vocalization Counts (CVC) accuracy.   Given the inaccuracy of far LENA tags, and in order to follow the LENA system procedure, we only counted vocalizations attributed to CHN and ignored those attributed to CHF. As shown in Figure (CVC), there is a strong association between clip-level counts estimated via the LENA system and those found in the human annotations: the Pearson correlation between the two was r = 0.71 (p = 0.00) when all clips were taken into account, and r = 0.77 (p = 0.00) when only clips with some child speech (i.e., excluding clips with 0 counts in both LENA and human annotations) were considered. This suggests that the LENA system captures differences in terms of number of child vocalizations across clips well.

However, users need more: They also interpret the absolute number of vocalizations found by LENA. Therefore, it is important to also bear in mind the absolute error rate and the relative error rate. The absolute error rate tells us, given a LENA estimate, how close the actual number may be. The relative error rate puts this number in relation to the actual

child voc counts (near only)

Figure 4. solid line corresponds to all clips, dashed clips corresponds to an analysis excluding clips where both the human and LENA said zero child vocalizations.

number of vocalizations tagged by the human coder. For instance, imagine that we find that LENA errs by 10 vocalizations according to the absolute error rate; this means that, on average across short clips like the ones used here, the numbers by LENA would be off by 10 vocalizations. We may think this number is small; by using the relative error rate, we can check whether it is small relative to the actual number: An error of 10 vocalizations would seem less problematic if there are 100 vocalizations on average (LENA would be just 10% off) than if there are 10 (LENA would be doubling the number of vocalizations).

The absolute error rate ranged from -47 to 30, with a mean of -1.63 and a median of 0. Since these numbers can be affected by silent clips, in which both humans and LENA agree

on there not being any vocalizations, we also calculated the absolute error rate excluding these clips. In this case, the absolute error rate ranged from -47 to 13, with a mean of -5.35 and a median of -4. As for relative error rates, these require the number in the denominator to be non-null. For this analysis, therefore, we need to remove the 460 clips in which the human annotator said there were no child vocalizations whatsoever. When we do this, the mean relative error rate ranged from -100 to 800, with a mean of -20.17 and a median of -44.44

Conversational turn counts (near only)

Figure 5. solid line corresponds to all clips, dashed clips corresponds to an analysis excluding clips where both the human and LENA said zero child turns.

Conversational Turn Counts (CTC) accuracy.   Again, we only considered "near" speaker categories in the turn count, and applied the same rule the LENA does,

where a turn can be from the key child to an adult or vice versa, and should happen within 5 seconds to be counted. The association between clip-level LENA and human CTC was weaker than that found for CVC: the Pearson correlation between the two was r = 0.55 (p = 0.00) when all clips were taken into account, and r = 0.47 (p = 0.00) when only clips with some child speech (i.e., excluding 322 clips with 0 counts in both LENA and human annotations) were considered. The absolute error rate ranged from -75 to 25, with a mean of -6.82 and a median of 0. The absolute error rate excluding clips where both human and LENA counts were zero ranged from -75 to 20, with a mean of -15.67 and a median of -12. As for relative error rates, these require the number in the denominator to be non-null. For this analysis, therefore, we need to remove the 427 clips in which the human annotator said there were no child-adult or adult-child turns whatsoever. When we do this, the mean relative error rate ranged from -100 to 366.67, with a mean of -64.64 and a median of -81.82

Adult Word Counts accuracy. One child in the SOD corpus was learning French. We have included this child to increase power, but results without this one child are nearly identical. The association between clip-level LENA and human AWC was strong: the Pearson correlation between the two was r=0.75 (p=0.00) when all clips were taken into account, and r=0.69 (p=0.00) when only clips with some child speech (i.e., excluding 309 clips with 0 counts in both LENA and human annotations) were considered. This suggests that the LENA system captures differences in terms of number of child vocalizations across clips well. The absolute error rate ranged from -211 to 157, with a mean of -0.20 and a median of 0. The absolute error rate excluding clips where both human and LENA counts were zero ranged from -211 to 157, with a mean of 0.46 and a median of -2. As for relative error rates, these require the number in the denominator to be non-null. For this analysis, therefore, we need to remove the 361 clips in which the human annotator said there were no child-adult or adult-child turns whatsoever. When we do this, the mean relative error rate ranged from -100 to 7400, with a mean of 55.04 and a median of -17.78

AWC

Figure 6. solid line corresponds to all clips, dashed clips corresponds to an analysis excluding clips where both the human and LENA said zero AWC.

Effects of age and differences across corpora. The preceding sections include results that are wholesale, over all corpora. However, we have reason to believe that performance could be higher for the corpora collected in North America (BER, WAR, SOD) than those collected in other English-speaking countries (L05) or non-English speaking populations (TSI). Additionally, our age ranges are wide, and in the case of TSI children, some of the children are older than the oldest children in the LENA training set. To assess whether accuracy varies as a function of corpora and child age, we fit mixed models as follows.

We predicted false alarm, miss, and confusion rates (when all "F" categories, TV, and

Figure 7. Diarization error rate as a function of corpus and child age (divided into quartiles). Only corpus-quartile combinations with more than 2 children are shown. Central tendency is median. Error bars indicate standard deviation over observations divided by the square root of the number of infants.

overlap were mapped onto silence, which yielded the best results in Section XX) from corpus, child age, and the interaction as fixed effects, child ID as random effect, on clips where there was some speech according to the human annotator. We followed up with an Analysis of Variance (type 2) to assess significance. In none of these analyses was corpus, child age, or their interaction significant.

For CVC, we fit a mixed model where CVC according to the human was predicted from CVC according to LENA, in interaction with corpus and age, as fixed factors; with child ID as random effect. An Analysis of Variance (type 2) found a triple interaction, suggesting that the predicted value of LENA with respect to human CVC depended on both the corpus and the child age; and a two-way interaction between CVC by LENA and corpus. To investigate these further, we fit a model where CVC according to the human was

predicted from CVC according to LENA in interaction with age (as fixed factors, with child ID as random) within each corpus separately. This revealed a significant interaction between LENA CVC and age for BER (indicating that the predictive value of LENA CVC increased with child age), whereas for the other four corpora this interaction was not significant, nor was the main effect of age, and only the LENA CVC emerged as a significant predictor of variance in child vocalization counts derived from human annotation.[1]

For CTC, we fit a mixed model where CTC according to the human was predicted from CTC according to LENA, in interaction with corpus and age, as fixed factors; with child ID as random effect. An Analysis of Variance (type 2) found a two-way interaction between CTC by LENA and corpus. To investigate this further, we fit the same regressions within each corpus separately.[2] These follow-up analyses revealed that CTC by LENA varied in its strength of prediction of human-tagged CTC across corpora (BER estimate = 1.16, SE of estimate = 0.16, t = 7.41; L05 (estimate = 1.62, SE of estimate = 0.21, t = 11.74) TSI estimate = 0.94, SE of estimate = 0.08, t = 11.74; SOD estimate = 0.96, SE of estimate = 0.22, t = 4.46; WAR estimate = 1.68, SE of estimate = 0.14, t = 12.22).

Discussion

The aim of the present study was to assess LENA accuracy across key outcome measures: speaker classification accuracy, adult word counts, child vocalization counts, and conversational turn counts. We did this using a method that avoided many of the limitations inherent in prior published analyses, which can lead to inflated accuracy rates; we included evaluations of all categories of input including noise, silence and overlap, and we used random and periodic sampling to select portions of the files for manual annotation, rather than relying on LENA's own output. In this way, we conducted an evaluation that is fully independent of the LENA algorithms automated output assessment, permitting a systematic, extensive, and independent evaluation of its key automated metrics. We also tested generalizability by analyzing LENA's performance across five different corpora; three based

on the same population, language and dialect that LENA was established for, and trained on (North American English), one that allowed us to test how accurately it captured a different dialect of English (UK English), and one that tested its performance in a totally different recording situation: a rural setting with large families and many children present, speaking a linguistically unrelated language.

Our first set of analyses tested LENA's overall accuracy, using established speech and talker segmentation metrics (false alarm rate, miss rate, confusion rate, and composite diarization error rate), and evaluated the pattern of errors in more detail, by assessing how LENA and human annotators agreed and disagreed (precision and recall). The overall diarization error rate was relatively high (79%), but this was mainly due to a high miss rate (missing or excluding speech that was there; 49%). The false alarm rate (identifying non-speech/silence as speech; 12%) and confusion rate (identifying voice type; 10%) were low.

However, the low overall confusion rate (10%) hid some considerable errors in parts of the system, since LENA performed much better with some talker categories than others. In terms of precision (to what extent do LENA tags contain what they say they contain), the system performed relatively well at identifying female voices (52% of frames tagged by LENA as FAN were coded as female adult by the human coders), and reasonably well at identifying the target child (39% of frames tagged by LENA as CHN were correct). However, the system performed substantially worse with other talker types (e.g. 8% and 6% for MAN and CXN, respectively); that is, less than a tenth of the frames that LENA tagged as being speech spoken by these speakers actually correspond to them.

In terms of recall (how accurately LENA captured the human annotations), performance for key child's speech was relatively robust; almost half of the frames tagged by humans being key child speech were captured by LENA under the CHN tag. However, recall was poorer for other talkers: only a third of adult female near speech (FAN) and less than 20% of adult male and other child speech were correctly tagged by LENA. How does LENA fare compared to other speech technology systems in terms of classification accuracy?

Although none have been tested on precisely the same data, there are indications that state-of-the-art speech technology can match LENA even without extensive training. For instance, Sell et al. 2017 obtained a diarization error rate of 60% (vs. 79% obtained here) on a different random sampling of the BER dataset from which the present work also draws (albeit focusing on clips that had some speech, which likely underestimate false alarm rates). This invites further interaction with the speech technology community to import such systems into our own analyses and studies.

Our second set of analyses tested the accuracy of three of LENA's aggregated counts; Child Vocalization Counts (CVC), Conversational Turn Counts (CTC) and Adult Word Counts (AWC). On one hand, we found strong associations between clip-level counts estimated via the LENA system and those from the human annotations for AWC and CVC, though performance was weaker for CTC. However, such correlational analyses do not establish whether LENA systematically over- or under-estimates. For this we examined absolute and relative error rate.

While the median error rate for CVC, CTC, and AWC was an encouraging 0%, this was likely due to many clips lacking vocalizations, turns, or adult words altogether. When such clips were excluded in the relative error analyses, median error rates rose substantially. Relative to humans, the LENA system missed roughly half of the relevant data for CVC, CTC, and AWC (-44%, -82%, and -18% relative error rates, respectively). Thus, LENA systematically underestimates the raw counts of its main quantitative measures - particularly child vocalizations and conversational turns, and to a lesser extent, adult words.

That said, LENA results were surprisingly robust to dialect, age, language and settings, as we found out in our final set of analyses, which tested how well all of these metrics generalize across corpora. We predicted that performance might be higher for the corpora collected in North America (BER, WAR, SOD) than those collected in other English-speaking countries (L05) or non-English speaking populations (TSI), and that accuracy might decrease for children older than those included in the LENA training set.

Contrary to our predictions, there were no significant differences in rates of false alarms, misses, or confusion by corpus, child age or their interaction, nor were there differences between North American and other corpora in terms of CVC or CTC. Instead, LENA's predictions varied across corpora in a way that could not be easily captured by age or dialect. For instance, LENA's CVC accuracy increased with age for BER whereas it seemed stable in the other corpora. However, note that here we are working with relatively small human-coded samples from each corpora; further work on bigger samples is required to verify these findings.

Overall, we conclude that there were very few language or dialect effects, and perhaps very few age effects. This is a promising finding for future speech technology solutions, since it suggests that diarization success may not require large quantities of highly specific training data. This is particularly important for the possibility of using LENA for languages unrelated to the (dialect of) English for which it was developed. That is, given the significant challenges in creating new datasets for training and testing speech technology, if it indeed turns out to be the case that automated speech processing (e.g. diarization) for daylong recordings is relatively language- or dialect-agnostic in its accuracy, it may be possible to use existing American English corpora with high hopes of generalizability.[3]

In sum, LENA performs relatively well in terms of overall accuracy, but there are pockets in the system (e.g. identifying male adult voices, establishing absolute numbers of child-adult turns) where the results of the algorithms are quite unreliable. Thus, whether LENA results are "good enough" for a given research, educational, or clinical study depends largely on the goals of each particular study. For example, above we have used terms such as "relatively/reasonably well" to describe precision rates of 52% (52% of frames tagged by LENA as FAN were coded as female adult by human coders) and 39% (39% of frames tagged as target child were also tagged as such by human coders). We do this partly because these rates are much higher than LENA's precision rates for other speakers (all below 10%) but also partly because our frame-based criteria is more stringent than many coding schemes

previously applied. However, that said, whether a particular accuracy rate can be considered "good enough", will depend on the purpose of the study. As a result, we conclude with a set of recommendations to help researchers make this determination for their goals.

What research goals can one pursue given LENA's performance?   In the present corpora, LENA's false alarm rate (i.e., identifying speech where there was none) was very low ($M = 12\%$). However its miss rate (missing speech that was actually there) was relatively high ($M = 49\%$). This makes it more suitable for studies in which it is extremely important not to "invent" speech that is not there but less suitable for studies in which capturing most, if not all, of the speech produced is crucial. Based on these findings, LENA would be a good tool for finding "high talk volume" parts of the day for a) careful further transcription (e.g. of low-frequency events like a certain grammatical construction of interest), b) annotation of specific speech characteristics (e.g. mean length of utterance), or c) comparing relative talk volume across samples. However, we advise caution in using LENA when raw quantity of speech is crucial for the research question, or when small differences in talk volume might have very significant theoretical consequences; this is often the case in clinical populations where children's own vocalizations can be an important diagnosis-relevant characteristic (e.g. in children who are deaf or hard of hearing, individuals with ASD, speech apraxia, etc.).

Similarly, although LENA's overall confusion rate (i.e. incorrectly identifying talkers, e.g. giving a "mother" tag for a "child" utterance) was very low (8%), this does not fully convey the level of accuracy for every talker type. In terms of precision, LENA's female adult and key child categorization was quite accurate, whereas precision was lower for male adults and other children, such that many of the frames labeled as male adult or other children did not in fact contain speech by these speaker types. In terms of recall, LENA was good at capturing speech by the key child as such, but recall was lower for the other voice categories, with very poor performance for speech by male adults and non-target children. We, thus, recommend caution before undertaking any analyses that rely on the accuracy

(precision and/or recall) of male adult and other children's speech. For example, if the goal is simply to calculate an overall adult word count (AWC), summing over male and female adult speakers, the fact that there is some confusion between MAN and FAN is likely not problematic. However, if the goal of the study is to compare the relative input from fathers and mothers, LENA tags are relatively unreliable and on our view, merit further manual vetting in most use cases. As another example, if the goal is to capture as many of the key child's vocalisations as possible, it might be worthwhile to pull out segments LENA labelled as non-target child (of which 22% was target child speech) as well, with human coders brought in to filter out non-target child speech.

However, while we recommend LENA users to be very careful in their use of LENA diarization and classification, especially for certain talker classes, our results for LENA count metrics suggest these derived counts may be accurate enough to serve well across a large variety of uses. To begin with, as far as it is possible to generalize from the limited range of samples tested here (children aged 2 to 58 months, learning North American English, UK English, or Tsimane) it seems that LENA performance does not vary a great deal across ages, dialects, language and home settings. Moreover, correlations between human and LENA clip-level counts were high to very high, suggesting that the software accurately captures differences in counts across clips (even when error rates were also high). These high correlations remained even when clips with counts equal to zero were removed from consideration, suggesting that LENA captures gradience in vocalization counts.

However, our finding that LENA generally underestimates the quantity of vocalization, turn or adult words deserves further consideration. While we feel caution is in order, further study is needed to fully understand the nature and extent of this limitation. Our clips were 1-2 minutes in length, and therefore they either tended to have very little speech or a lot of it. Error rates over hours could be smaller, because local errors average out; or greater; if the LENA system systematically underestimate counts. In a LENA technical report, AWC accuracy was variable across two 12-hour recordings: 1% lower than human transcription for

one child, but 27% lower for a second child. This same report notes that AWC accuracy quickly plateaus as recording time increases beyond one hour, leveling to 5-10% in recording >2 hours [LTR-05-2].

Thus, it is important for further work to help establish the systematicity in LENA's estimates: if underestimates are robust and systematic, it may be possible to develop a correlation factor to account for this bias. However, this bias may be challenging to nail down precisely. For instance, a recent study in Finland documented an OVERestimate by LENA, suggested that this raw bias may be context specific [Elo dissertation ref]. Nonetheless, we are overall hopeful that the reliability metrics we provide here will be relevant for researchers working with different populations. We strongly encourage full reporting of LENA validation to bolster this literature for the whole community.

How to test the reliability of LENA's results. Although we did not find large differences across languages, ages, dialects and settings, we recommend that researchers test the reliability of LENA counts in their own samples, especially if they are collecting data from families living in different environments from those assessed here. Below, we provide some guidelines for how to go about this. Note that this requires downloading the audio (.wav) file as well as the LENA output file. However, the LENA recorder itself produces good quality audio output, so we recommend that researchers always download the audio file in any case, so that it is available for future analysis. First, we recommend a literature search, to determine whether a similar sample has been studied in the past for which there exists reliability data (see for example, Cristia et al, submitted, for a systematic review). If no studies exist, draw 10 x 2 minutes randomly from 10 children. This is about 3h20min of data, which takes roughly 60h to annotate in our experience. We recommend training annotators using ACLEW Annotation Scheme https://osf.io/b2jep/ and then using DiViMe (divime.readthedocs.io, Le Franc et al 2018 Interspeech) to estimate the accuracy for the sample. Extract the classification accuracy measures used here (% misses, % false alarms, % confusions) from the annotated data and sum it to provide a total diarization error rate

using the recipes provided in that package. The supplementary materials to the present paper contain scripts that allow users to extract CVC, CTC, and AWC from LENA and annotations made using AAS. Separately, estimate the accuracy needed for the study. For instance, suppose we have an evaluation of an intervention where we expect treatment children to hear 20% more speech than controls, or an individual differences study where we expect that the lower 5th of the children hear 20% less speech than the top 5th. If the intended measure used to compare groups has an error rate larger than the effect predicted (e.g. the CVC error rate we find here), a different algorithm or outcome metric is recommended (see DiViMe for some alternative algorithms).

Conclusions.   In conclusion, in this study, we have provided a broad evaluation of LENA accuracy across its key outcome measures (classification accuracy, adult word counts, child vocalization counts, and conversational turn counts), and its generalizability across different dialects, languages, ages and settings. We have provided some recommendations for how to use LENA in future studies most effectively, and how to test the accuracy of the LENA algorithms on particular samples of data.

There are, however, a number of areas of research that we have not addressed. For example, we have not investigated how accurately LENA detects individual variation across children or families. It would be particularly useful to know whether LENA can classify children with the sensitivity and specificity needed for accurate identification of language disorders. Oller et al, 2010 used LENA to differentiate vocalizations from 232 typically developing children and children with autism or language delay with a high degree of accuracy. However, key to this was the use of additional algorithms, not yet available with LENA, to identify and classify the acoustic features of "speech-related vocal islands" (SVIs). Further work is, thus, needed here.

Even if it turns out that LENA is not accurate enough to classify children precisely, it may be accurate enough to capture the rank order of individual children's language growth, which can provide useful information about the relative language level of children in a

sample or population (see e.g. Gilkerson et al, 2017)). Similarly, LENA may be able to track a child's development relatively accurately over time;iIt may not capture accurately the precise number of child vocalisations produced over time, but it may track developmental trajectory (e.g. the slope of growth) relatively well. Finally, although our results suggest that aspects of LENA's output may be relatively robust to differences across languages/dialects, we need more evidence of how it fares when tracking the language, and language environment, of multilingual children in multilingual homes (see Orena et al (date??) for some evidence that LENA is reliable in French-English bilingual environments although, as in the present paper, it underestimated adult word count). More work is needed to investigate these research questions.

Acknowledgments

References